17 Co-word and Keyword Co-occurrence

R. Heller · Date: 2026-05-09 · Last changed: 2026-10-31

17.1 Learning objectives

After completing this chapter, you will be able to:

- Explain the principles of co-word analysis and its role in science mapping
- Extract and clean keywords from OpenAlex data (author keywords and concepts)
- Build a keyword co-occurrence matrix and convert it to a network
- Apply community detection to identify topical clusters
- Visualise keyword co-occurrence networks with interpretable layouts

17.2 Setup

```
library(tidyverse)
library(openalexR)
library(igraph)
library(tidygraph)
library(ggraph)
library(glue)
library(gt)

set.seed(20260509)

source(here::here("R", "api_helpers.R"))
source(here::here("R", "utils.R"))
source(here::here("R", "sci_palette.R"))
```

17.3 Conceptual background

Co-word analysis maps the conceptual structure of a research field by examining which terms appear together in the same documents. Introduced by Callon et al. (1983), the method assumes that when two keywords repeatedly co-occur across publications, they are thematically related. The resulting co-occurrence network reveals the topical landscape of a field: clusters of densely connected keywords represent coherent research themes, while bridges between clusters indicate interdisciplinary connections.

The method works with several types of terms:

- **Author keywords**: Terms selected by the paper's authors. These are intentional descriptors but suffer from inconsistency — authors may use different terms for the same concept ("machine learning" vs. "ML" vs. "statistical learning").
- **Indexed keywords**: Terms assigned by database indexers (e.g., MeSH terms in PubMed). More consistent but only available in some databases.

- **OpenAlex concepts**: Algorithmically assigned topics at multiple levels of a hierarchical taxonomy. These provide broad coverage but may miss nuanced distinctions (Priem et al. 2022).
- **Title/abstract words**: Extracted via text mining. Comprehensive but noisy; requires extensive preprocessing.

Normalisation is important for co-word networks. Raw co-occurrence counts favour high-frequency terms. Common normalisations include the **association strength** (equivalent to pointwise mutual information) and the **equivalence index** (cosine of the co-occurrence vector). Waltman et al. (2010) demonstrated that association strength produces the most balanced network structures for bibliometric mapping.

Co-word analysis complements citation-based methods (16.3). Citation networks reveal intellectual influence; co-word networks reveal thematic content. Combining both provides a more complete picture of a field's structure.

17.4 Worked example

17.4.1 Extracting keywords

We use OpenAlex-assigned topic labels as terms in a sample of scientometrics research. These are algorithmic labels, not author-supplied keywords. The flattened topic table also includes parent subfields, fields, and domains; retain only topic identifiers (`T...`) and count each term once per work.

```
works ← fetch_openalex(
  entity = "works",
  primary_location.source.id = "S148561398",
  from_publication_date = "2019-01-01",
  to_publication_date = "2023-12-31",
  options = list(sample = 500, seed = 42)
)
```

```
keywords ← works ▷
  select(id, topics) ▷
  unnest(topics, names_sep = "_") ▷
  filter(str_detect(topics_id, "(^|/)T[0-9]+$")) ▷
  filter(topics_display_name ≠ "") ▷
  select(work_id = id, keyword = topics_display_name) ▷
  mutate(keyword = str_to_lower(str_trim(keyword))) ▷
  filter(!is.na(keyword), nchar(keyword) ⩾ 3) ▷
  distinct(work_id, keyword)

kw_counts ← keywords ▷
  count(keyword, sort = TRUE)

cat(glue("Total keyword occurrences: {nrow(keywords)}\n"))
```

```
#> Total keyword occurrences: 1497
```

```r
cat(glue("Unique keywords: {n_distinct(keywords$keyword)}\n"))
```

```
#> Unique keywords: 135
```

```r
kw_counts |>
  head(20) |>
  mutate(keyword = fct_reorder(keyword, n)) |>
  ggplot(aes(x = n, y = keyword)) +
  geom_col(fill = palette_sci(1)) +
  labs(x = "Frequency", y = NULL) +
  theme_sci()
```

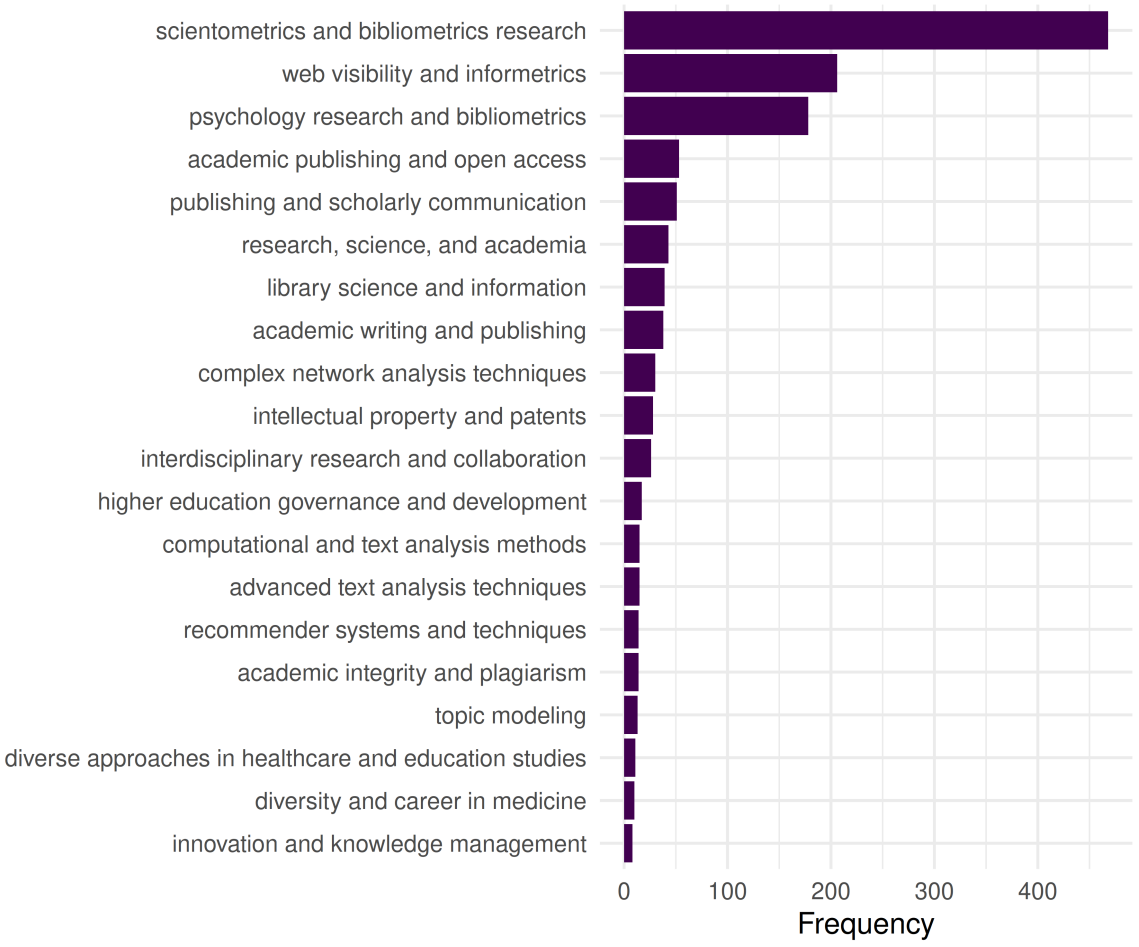


Figure 17.1: Top 20 most frequent assigned topics used as terms in the Scientometrics sample.

17.4.2 Building the co-occurrence network

We create edges between keywords that appear in the same paper.

```
kw_frequent ← kw_counts ▷
  filter(n ⩾ 5) ▷
  pull(keyword)

kw_filtered ← keywords ▷
  filter(keyword %in% kw_frequent)

coword_pairs ← kw_filtered ▷
  inner_join(kw_filtered, by = "work_id", suffix = c("_a", "_b"),
             relationship = "many-to-many") ▷
  filter(keyword_a < keyword_b) ▷
  count(keyword_a, keyword_b, name = "cooccurrence")

coword_top ← coword_pairs ▷
  filter(cooccurrence ⩾ 3)

g_kw ← graph_from_data_frame(
  coword_top ▷ select(keyword_a, keyword_b, weight = cooccurrence),
  directed = FALSE
) ▷
  simplify(edge.attr.comb = list(weight = "sum"))

cat(glue("Keyword network: {vcount(g_kw)} nodes, {ecount(g_kw)} edges\n"))
```

```
#> Keyword network: 33 nodes, 68 edges
```

17.4.3 Community detection for topical clusters

```
V(g_kw)$degree ← degree(g_kw)
V(g_kw)$strength ← strength(g_kw)

communities ← cluster_leiden(g_kw, resolution_parameter = 0.8,
                             objective_function = "modularity")
V(g_kw)$community ← as.factor(membership(communities))

cat(glue("Communities: {length(unique(membership(communities)))}\n"))
```

```
#> Communities: 3
```

```
cat(glue("Modularity: {round(modularity(g_kw, membership(communities)), 3)}\n"))
```

```
#> Modularity: 0.036
```

```r
community_summary ← tibble(
  keyword = V(g_kw)$name,
  community = V(g_kw)$community,
  strength = V(g_kw)$strength
) ▷
  group_by(community) ▷
  slice_max(strength, n = 5) ▷
  summarise(top_keywords = paste(keyword, collapse = ", "),
            n_keywords = n(), .groups = "drop")

community_summary ▷ gt()
```

community	top_keywords	n_keywords
1	scientometrics and bibliometrics research, web visibility and informetrics, psychology research and bibliometrics, academic publishing and open access, publishing and scholarly communication	5
2	intellectual property and patents, advanced text analysis techniques, computational and text analysis methods, topic modeling, innovation and knowledge management	5
3	recommender systems and techniques, information retrieval and search behavior	2

17.4.4 Visualisation

```r
set.seed(42)
layout ← create_layout(as_tbl_graph(g_kw), layout = "fr")

ggraph(layout) +
  geom_edge_link(aes(width = weight), alpha = 0.15, colour = "grey60") +
  scale_edge_width_continuous(range = c(0.3, 2), guide = "none") +
  geom_node_point(aes(size = degree, colour = community), alpha = 0.8) +
  geom_node_text(aes(label = ifelse(degree > quantile(degree, 0.85),
                                    name, NA_character_)),
                 repel = TRUE, size = 2.5, max.overlaps = 20, na.rm = TRUE) +
  scale_size_continuous(range = c(1, 6), guide = "none") +
  scale_colour_manual(values = palette_sci(
    n_distinct(V(g_kw)$community)
  )) +
  labs(colour = "Cluster") +
  theme_void(base_family = "sans", base_size = 11) +
  theme(legend.position = "bottom")
```

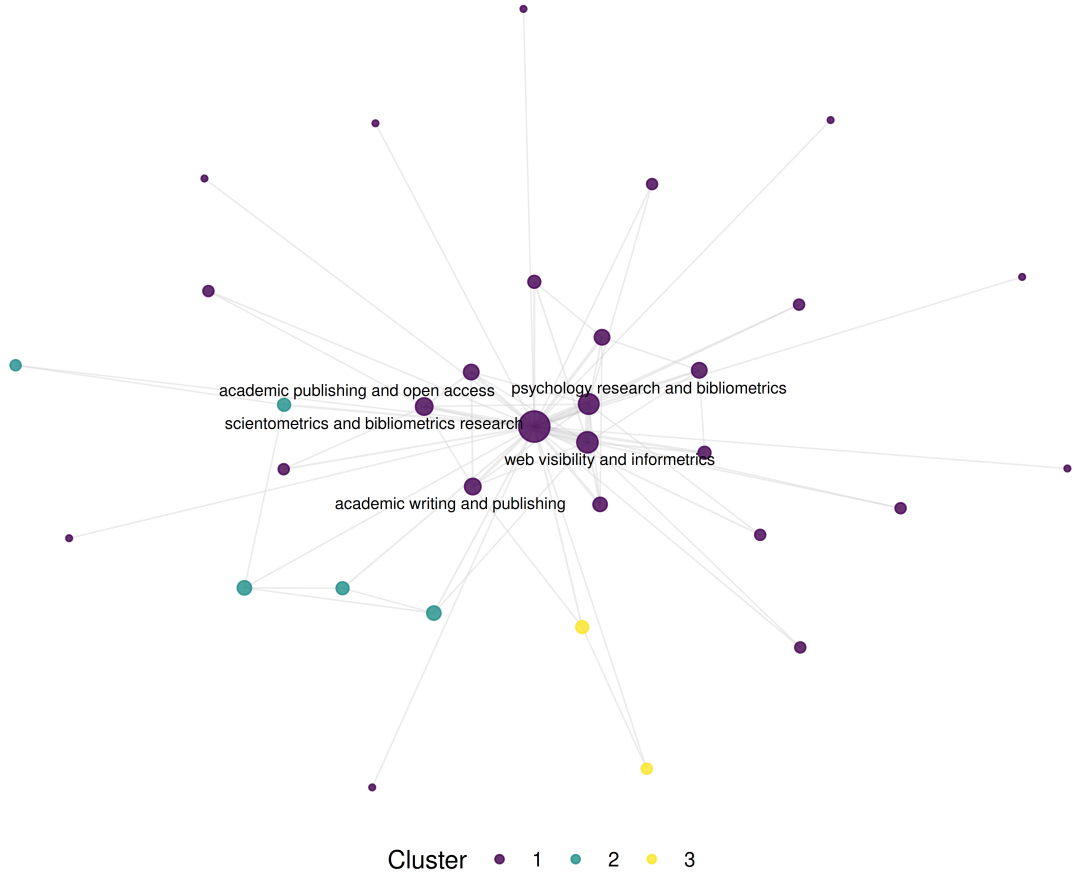


Figure 17.2: Keyword co-occurrence network coloured by topical cluster.

17.5 Diagnostics and interpretation

- **Keyword cleaning**: Inconsistent terminology inflates the vocabulary and creates spurious nodes. Standardise spelling, merge synonyms (e.g., “h-index” and “hirsch index”), and remove generic terms (“research”, “analysis”) that co-occur with everything but convey no topical information.
- **Frequency threshold**: Including rare keywords produces large, sparse, unreadable networks. Start with keywords appearing in at least 5 papers and adjust based on corpus size.
- **Community interpretability**: Each community should correspond to a recognisable research theme. If communities are uninterpretable, the resolution parameter may need adjustment or keywords need further cleaning.
- **Centrality interpretation**: High-degree keywords are thematically central (used across many contexts). High-betweenness keywords bridge distinct topics and may represent interdisciplinary concepts.

17.6 Limitations and responsible use

17.7 Limitations and responsible use

- **Vocabulary inconsistency.** Author keywords are not controlled vocabulary. The same concept may appear under multiple terms, fragmenting what should be a single node. Merging synonyms requires domain expertise.
- **Algorithmic concepts.** OpenAlex concepts are assigned by machine learning and may contain errors — especially at fine-grained levels. Always validate topic assignments by reading sample papers (Priem et al. 2022).
- **Static snapshots.** A co-word network represents a time-averaged view. Emerging topics with few publications may be invisible. Consider building temporal slices to track field evolution.
- **Not evaluative.** Popular keywords indicate research activity, not research quality or societal impact. Do not equate topical centrality with importance (Hicks et al. 2015).

17.8 Common pitfalls

17.9 Common pitfalls

- **Not cleaning keywords.** Punctuation variants ("co-authorship" vs. "coauthorship"), case differences, and trailing whitespace create duplicate nodes. Clean before building the network.
- **Including stop-keywords.** Generic terms like "research", "study", "analysis" co-occur with everything and dominate the network without conveying topical information. Remove them.
- **Comparing co-occurrence counts across corpora of different sizes.** Larger corpora produce higher co-occurrence counts mechanically. Normalise or use relative measures.
- **Over-reading small clusters.** A cluster with two or three keywords may reflect a single paper, not a research theme. Report cluster size alongside content.

17.10 Exercises

1. **Author keywords vs. concepts.** Build co-occurrence networks from both author keywords (if available) and OpenAlex concepts for the same corpus. Compare the resulting topical clusters.

2. **Temporal evolution.** Split the corpus into yearly slices. For each year, build a co-word network and identify the top 5 keywords by degree. Which keywords are consistently central, and which emerge over time?
3. **Strategic diagram.** Compute density (internal cohesion) and centrality (external connections) for each keyword cluster. Plot them on a strategic diagram (density vs. centrality quadrants). Which clusters are core themes and which are peripheral?
4. **Normalisation comparison.** Build two networks from the same data: one with raw co-occurrence counts and one with association strength normalisation. How do the most central keywords differ?

17.11 Solutions

Solutions are provided in E.

17.12 Further reading

- Callon et al. (1983) — The foundational paper on co-word analysis for mapping science.
- Waltman et al. (2010) — Unified framework for bibliometric network construction, including keyword networks.
- Aria and Cuccurullo (2017) — `bibliometrix` includes co-word analysis via the `biblioNetwork()` function.
- Priem et al. (2022) — OpenAlex concepts and topic classification.

17.13 Session info

```
sessionInfo()
```

```
#> R version 4.6.1 (2026-06-24)
#> Platform: x86_64-pc-linux-gnu
#> Running under: Ubuntu 24.04.5 LTS
#>
#> Matrix products: default
#> BLAS:   /usr/lib/x86_64-linux-gnu/openblas-pthread/libblas.so.3
#> LAPACK: /usr/lib/x86_64-linux-gnu/openblas-pthread/libopenblasp-r0.3.26.so;  LAPACK
version 3.12.0
#>
#> locale:
#>  [1] LC_CTYPE=C.UTF-8       LC_NUMERIC=C           LC_TIME=C.UTF-8
LC_COLLATE=C.UTF-8     LC_MONETARY=C.UTF-8
#>  [6] LC_MESSAGES=C.UTF-8    LC_PAPER=C.UTF-8       LC_NAME=C
LC_ADDRESS=C           LC_TELEPHONE=C
#> [11] LC_MEASUREMENT=C.UTF-8 LC_IDENTIFICATION=C
#>
#> time zone: UTC
#> tzcode source: system (glibc)
#>
#> attached base packages:
#> [1] stats     graphics  grDevices utils     datasets  methods   base
#>
#> other attached packages:
#>  [1] ggraph_2.2.2       tidygraph_1.3.1    igraph_2.3.4       quanteda_4.5.1
pdftools_3.9.1     arrow_26.0.0
#>  [7] bibliometrix_5.5.0 RefManageR_1.4.0   bib2df_1.1.2.0     rcrossref_1.2.1
gt_1.4.0           tidytext_0.4.3
#> [13] glue_1.8.1         openalexR_3.1.2    lubridate_1.9.5    forcats_1.0.1
stringr_1.6.0      dplyr_1.2.1
#> [19] purrr_1.2.2        readr_2.2.0        tidyr_1.3.2        tibble_3.3.1
ggplot2_4.0.3      tidyverse_2.0.0
#>
#> loaded via a namespace (and not attached):
#>   [1] bibtex_0.5.3           RColorBrewer_1.1-3     rstudioapi_0.19.0
jsonlite_2.0.0         magrittr_2.0.5
#>   [6] farver_2.1.2           rmarkdown_2.32         fs_2.1.0
vctrs_0.7.3            memoise_2.0.1
#>  [11] askpass_1.2.1          base64enc_0.1-6        htmltools_0.5.9
contentanalysis_1.1.1  curl_8.1.0
#>  [16] janeaustenr_1.0.0      cellranger_1.1.0       sass_0.4.10
bslib_0.12.0           htmlwidgets_1.6.4
#>  [21] tokenizers_0.3.0       plyr_1.8.9             httr2_1.3.0
plotly_4.12.1          cachem_1.1.0
#>  [26] dimensionsR_0.0.3      mime_0.13              lifecycle_1.0.5
pkgconfig_2.0.3        Matrix_1.7-5
#>  [31] R6_2.6.1               fastmap_1.2.0          shiny_1.14.0
digest_0.6.39          shinycssloaders_1.1.0
#>  [36] rprojroot_2.1.1        SnowballC_0.7.1        labeling_0.4.3
urltools_1.7.3.1       timechange_0.4.0
#>  [41] polyclip_1.10-7        httr_1.4.9             compiler_4.6.1
here_1.0.2             bit64_4.8.6
#>  [46] withr_3.0.3            S7_0.2.2               viridis_0.6.5
ggforce_0.5.0          MASS_7.3-65
#>  [51] bibliometrixData_0.3.0 tools_4.6.1            otel_0.2.0
stopwords_2.3          zip_3.0.2
#>  [56] httpuv_1.6.17          rentrez_1.2.4          promises_1.5.0
grid_4.6.1             stringdist_0.9.17
#>  [61] generics_0.1.4         gtable_0.3.6           tzdb_0.5.0
rscopus_0.9.0          ca_0.72
#>  [66] data.table_1.18.6.1    hms_1.1.4              xml2_1.6.0
utf8_1.2.6             ggrepel_0.9.8
#>  [71] pillar_1.11.1          later_1.4.8            tweenr_2.0.3
lattice_0.22-9         bit_4.6.0
```

```
#>  [76] tidyselect_1.2.1       miniUI_0.1.2           downlit_0.4.5
knitr_1.52             gridExtra_2.3.1
#>  [81] bookdown_0.48          crul_1.6.0             xfun_0.62
graphlayouts_1.2.5     DT_0.34.0
#>  [86] humaniformat_0.6.0     visNetwork_2.1.4       stringi_1.8.9
qpdf_1.4.1             yaml_2.3.12
#>  [91] evaluate_1.0.5         httpcode_0.3.0         cli_3.6.6
xtable_1.8-8           jquerylib_0.1.4
#>  [96] Rcpp_1.1.2             readxl_1.6.0           triebeard_0.4.1
XML_3.99-0.25          parallel_4.6.1
#> [101] assertthat_0.2.1       pubmedR_1.0.2          viridisLite_0.4.3
scales_1.4.0           openxlsx_4.2.9
#> [106] rlang_1.3.0            fastmatch_1.1-8
```