23 Keyword Extraction

23.1 Learning objectives

After completing this chapter, you will be able to:

- Explain the RAKE and TextRank keyword extraction algorithms
- Apply keyword extraction to bibliometric abstracts using R packages
- Compare automatically extracted keywords with author-assigned keywords
- Evaluate extraction quality using precision, recall, and F1 metrics
- Decide when automatic keyword extraction adds value over existing metadata

23.2 Setup

```
library(tidyverse)
library(openalexR)
library(quanteda)
library(quanteda.textstats)
library(glue)
library(gt)

set.seed(20260509)

source(here::here("R", "api_helpers.R"))
source(here::here("R", "rake_keywords.R"))
source(here::here("R", "utils.R"))
source(here::here("R", "sci_palette.R"))
```

23.3 Conceptual background

Keywords serve as compact descriptors of a document's content. In bibliometrics, they are used for co-word analysis (17.3), topic mapping, and search. Three sources of keywords exist:

1. **Author keywords**: Selected by the paper's authors. Intentional but inconsistent — different authors use different terms for the same concept.
2. **Indexed terms**: Assigned by database curators (MeSH in PubMed, descriptors in WoS). Consistent but labour-intensive and available only in some databases.
3. **Extracted keywords**: Automatically derived from titles and abstracts using algorithms. Scalable and consistent but may miss domain nuance.

RAKE (Rapid Automatic Keyword Extraction) identifies candidate keywords by finding sequences of content words delimited by stopwords or punctuation. It scores candidates by the ratio of word weighted degree (including the word itself within candidate phrases) to word frequency, favouring multi-word phrases that appear in specific contexts. RAKE is fast and unsupervised but does not consider document-level or corpus-level statistics.

TextRank adapts the PageRank algorithm to text: words are nodes in a graph, edges connect co-occurring words within a sliding window, and importance is computed iteratively. High-ranking words are selected as keywords. TextRank captures contextual importance better than frequency-based methods but is more computationally expensive.

Both methods extract keywords from individual documents. **TF-IDF** (21.3) provides a complementary corpus-level view, identifying terms that distinguish a document from the broader collection.

Evaluating keyword extraction requires ground truth. Author keywords are the most common benchmark, but they are imperfect: authors may omit obvious keywords or include overly specific terms. Agreement depends on the corpus, matching rules, number of keywords, and reference annotation. Define these choices before evaluation and report observed precision and recall; no universal acceptance percentage is assumed here.

23.4 Worked example

23.4.1 Fetching data with keywords

```
works ← fetch_openalex(
  entity = "works",
  primary_location.source.id = "S148561398",
  from_publication_date = "2021-01-01",
  to_publication_date = "2023-12-31",
  options = list(sample = 200, seed = 42)
)

text_df ← works ▷
  filter(!is.na(abstract), nchar(abstract) > 100) ▷
  transmute(
    doc_id = id,
    title = display_name,
    abstract = abstract,
    text = paste(display_name, abstract, sep = ". ")
  )

cat(glue("Documents with abstracts: {nrow(text_df)}\n"))
```

```
#> Documents with abstracts: 67
```

23.4.2 RAKE keyword extraction

```
rake_extract ← function(text, n_keywords = 10) {
  rake_keywords(text, stopwords = quanteda::stopwords("en"),
                n_keywords = n_keywords)$keyword
}

text_df ← text_df ▷
  mutate(rake_keywords = map(text, \(t) rake_extract(t, n_keywords = 8)))

text_df ▷
  head(3) ▷
  select(title, rake_keywords) ▷
  mutate(rake_keywords = map_chr(rake_keywords, paste, collapse = "; ")) ▷
  gt()
```

title	rake_keywords
Market forces in Italian academia today (and yesterday)	better scholars travel longer distances; consistent nesting procedure; greater expected utility; choosing private universities; italian academia today; mapping current scholars; current academic market; sorting effect delineates
Why do papers from international collaborations get more citations? A bibliometric analysis of Library and Information Science papers	achieve relevant research outputs; countries must therefore consider; science mapping analysis approach; achieve greater scientific impact; abstract scientific activity; become increasingly complex; thus becomes necessary; information science category
Research on Covid-19: a disruptive phenomenon for bibliometrics	present study also confirms; san francisco declaration; world recent history; present study shows; highest disruptive event; bibliometrics provide methods; journal impact factor; worldwide academic research

23.4.3 TF-IDF based keyword extraction

```
corp ← corpus(text_df, docid_field = "doc_id", text_field = "text")

toks ← tokens(corp, remove_punct = TRUE, remove_numbers = TRUE) ▷
  tokens_tolower() ▷
  tokens_remove(stopwords("en")) ▷
  tokens_remove(c("study", "paper", "results", "research", "analysis"))

dfmat ← dfm(toks) ▷
  dfm_trim(min_termfreq = 2)
dfmat_tfidf ← dfm_tfidf(dfmat)

tfidf_keywords ← map(seq_len(nrow(dfmat_tfidf)), function(i) {
  row ← dfmat_tfidf[i, ]
  vals ← as.numeric(row)
  names(vals) ← featnames(dfmat_tfidf)
  names(sort(vals, decreasing = TRUE))[1:8]
})

text_df$tfidf_keywords ← tfidf_keywords
```

23.4.4 Corpus-level keyword frequency

```
all_kw ← text_df ▷
  pull(rake_keywords) ▷
  unlist() ▷
  tibble(keyword = _) ▷
  count(keyword, sort = TRUE)

all_kw ▷
  head(20) ▷
  mutate(keyword = fct_reorder(keyword, n)) ▷
  ggplot(aes(x = n, y = keyword)) +
  geom_col(fill = palette_sci(1)) +
  labs(x = "Extraction frequency", y = NULL) +
  theme_sci()
```

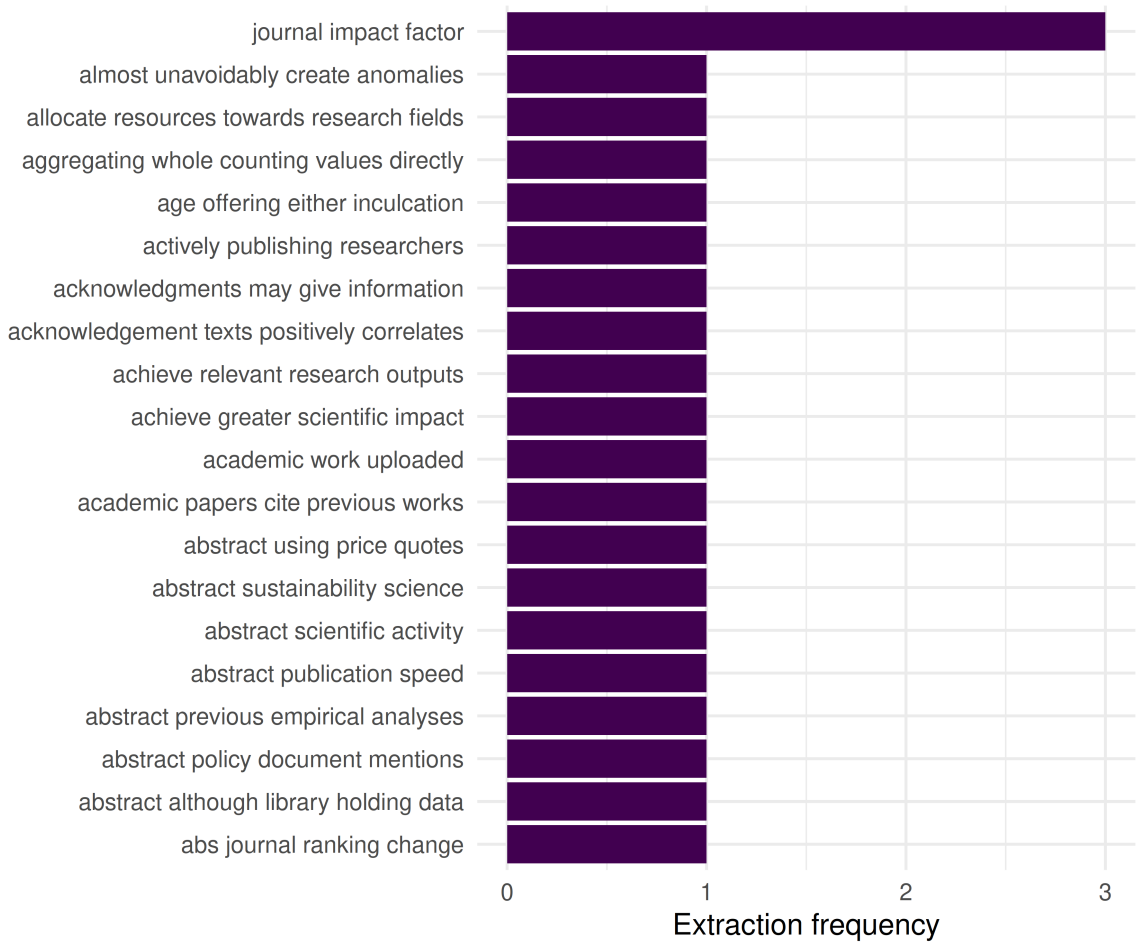


Figure 23.1: Top 20 keywords by total extraction frequency across the corpus.

23.4.5 Comparing extraction methods

```
overlap ← map2_dbl(text_df$rake_keywords, text_df$tfidf_keywords,
                   \(r, t) length(intersect(r, t)) / length(union(r, t)))

cat(glue("Mean Jaccard overlap (RAKE vs TF-IDF): {round(mean(overlap, na.rm = TRUE),
3)}\n"))
```

```
#> Mean Jaccard overlap (RAKE vs TF-IDF): 0.001
```

```
tibble(jaccard = overlap) ▷
  ggplot(aes(x = jaccard)) +
  geom_histogram(binwidth = 0.05, fill = palette_sci(1), colour = "white") +
  labs(x = "Jaccard similarity", y = "Documents") +
  theme_sci()
```

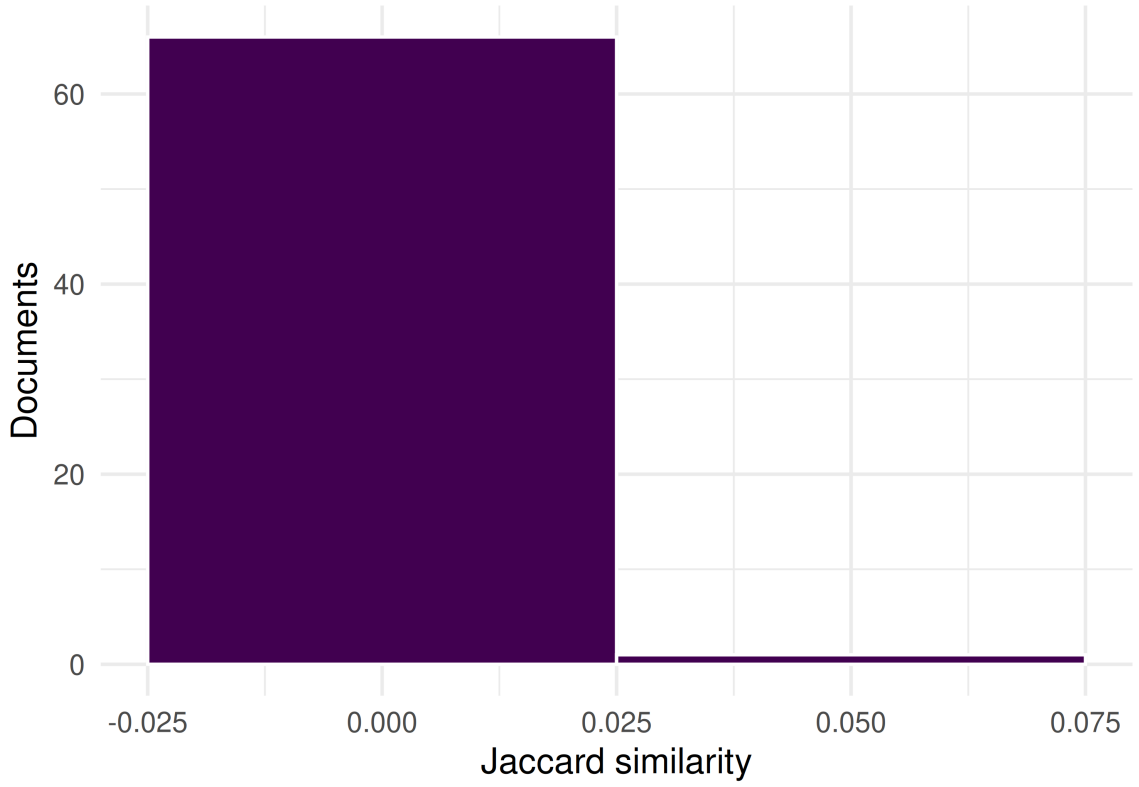


Figure 23.2: Distribution of Jaccard similarity between RAKE and TF-IDF keyword sets.

23.5 Diagnostics and interpretation

- **Extraction quality**: Compare extracted keywords against author keywords for a random sample of 20–30 papers. Specify exact or normalized matching, report precision and recall, and inspect disagreements rather than applying an unsupported universal threshold.
- **Redundancy**: Check for near-duplicate keywords (singular/plural, abbreviation/full form). Post-processing to merge variants improves quality.
- **Domain specificity**: Generic terms (“model”, “method”, “data”) are often extracted but carry no topical information. Consider a domain-specific stopword list.
- **Coverage**: Some documents yield few meaningful keywords, especially short abstracts. Flag documents where fewer than 3 keywords were extracted.

23.6 Limitations and responsible use

23.7 Limitations and responsible use

- **No method is comprehensive.** Keyword extraction captures salient terms but misses concepts expressed through circumlocution or implicit reference. Human keywords remain valuable.
- **Language dependence.** RAKE and TextRank work best for English. Stopword lists and tokenisation rules may fail for other languages.
- **Evaluation is subjective.** There is no single correct keyword set for a document. Low agreement between methods or between extracted and author keywords does not necessarily indicate failure.
- **Not a replacement for indexing.** Professional indexing (e.g., MeSH) follows controlled vocabularies and domain expertise. Automated extraction is a complement, not a substitute (Hicks et al. 2015).

23.8 Common pitfalls

23.9 Common pitfalls

- **Using raw frequencies as keywords.** The most frequent words in a document are often stopwords or near-stopwords. TF-IDF or RAKE scoring is essential.
- **Not normalising keywords.** “bibliometric” and “bibliometrics” should be the same keyword. Apply stemming or lemmatisation before evaluation.
- **Evaluating on titles only.** Titles are too short for reliable keyword extraction. Use abstracts or title + abstract combined.
- **Ignoring multi-word terms.** Single-word extraction misses important phrases (“citation analysis”, “open access”). Use n-grams or phrase detection.

23.10 Exercises

1. **Precision and recall.** For 20 papers with author keywords, compute precision and recall of RAKE-extracted keywords against author keywords. Report the mean and range.

2. **Bigram keywords.** Extract bigrams from abstracts and score them by TF-IDF. What multi-word terms emerge that single-word extraction misses?

3. **Keyword enrichment.** For papers without author keywords, use extracted keywords to build a co-word network. Does the resulting topical structure resemble that from author keywords?

23.11 Solutions

Solutions are provided in E.

23.12 Further reading

- Silge and Robinson (2017) — Keyword and term extraction in R with tidy tools.
- Aria and Cuccurullo (2017) — `bibliometrix` keyword analysis capabilities.
- Callon et al. (1983) — Co-word analysis, the downstream use case for extracted keywords.
- Priem et al. (2022) — OpenAlex concept tagging as an alternative to keyword extraction.

23.13 Session info

```
sessionInfo()
```

```
#> R version 4.6.1 (2026-06-24)
#> Platform: x86_64-pc-linux-gnu
#> Running under: Ubuntu 24.04.5 LTS
#>
#> Matrix products: default
#> BLAS:   /usr/lib/x86_64-linux-gnu/openblas-pthread/libblas.so.3
#> LAPACK: /usr/lib/x86_64-linux-gnu/openblas-pthread/libopenblasp-r0.3.26.so;  LAPACK
version 3.12.0
#>
#> locale:
#>  [1] LC_CTYPE=C.UTF-8       LC_NUMERIC=C           LC_TIME=C.UTF-8
#>  [4] LC_COLLATE=C.UTF-8     LC_MONETARY=C.UTF-8    LC_MESSAGES=C.UTF-8
#>  [7] LC_PAPER=C.UTF-8       LC_NAME=C              LC_ADDRESS=C
#> [10] LC_TELEPHONE=C         LC_MEASUREMENT=C.UTF-8 LC_IDENTIFICATION=C
#>
#> time zone: UTC
#> tzcode source: system (glibc)
#>
#> attached base packages:
#> [1] stats     graphics  grDevices utils     datasets  methods   base
#>
#> other attached packages:
#>  [1] stm_1.3.8                 topicmodels_0.2-17
#>  [3] quanteda.textstats_0.97.2 visNetwork_2.1.4
#>  [5] ggraph_2.2.2              tidygraph_1.3.1
#>  [7] igraph_2.3.4              quanteda_4.5.1
#>  [9] pdftools_3.9.1            arrow_26.0.0
#> [11] bibliometrix_5.5.0        RefManageR_1.4.0
#> [13] bib2df_1.1.2.0            rcrossref_1.2.1
#> [15] gt_1.4.0                  tidytext_0.4.3
#> [17] glue_1.8.1                openalexR_3.1.2
#> [19] lubridate_1.9.5           forcats_1.0.1
#> [21] stringr_1.6.0             dplyr_1.2.1
#> [23] purrr_1.2.2               readr_2.2.0
#> [25] tidyr_1.3.2               tibble_3.3.1
#> [27] ggplot2_4.0.3             tidyverse_2.0.0
#>
#> loaded via a namespace (and not attached):
#>   [1] bibtex_0.5.3           RColorBrewer_1.1-3     rstudioapi_0.19.0
#>   [4] jsonlite_2.0.0         magrittr_2.0.5         modeltools_0.2-25
#>   [7] farver_2.1.2           rmarkdown_2.32         fs_2.1.0
#>  [10] vctrs_0.7.3            memoise_2.0.1          askpass_1.2.1
#>  [13] base64enc_0.1-6        htmltools_0.5.9        contentanalysis_1.1.1
#>  [16] curl_8.1.0             janeaustenr_1.0.0      cellranger_1.1.0
#>  [19] sass_0.4.10            bslib_0.12.0           htmlwidgets_1.6.4
#>  [22] tokenizers_0.3.0       plyr_1.8.9             httr2_1.3.0
#>  [25] plotly_4.12.1          cachem_1.1.0           dimensionsR_0.0.3
#>  [28] mime_0.13              lifecycle_1.0.5        pkgconfig_2.0.3
#>  [31] Matrix_1.7-5           R6_2.6.1               fastmap_1.2.0
#>  [34] shiny_1.14.0           digest_0.6.39          patchwork_1.3.2
#>  [37] shinycssloaders_1.1.0  rprojroot_2.1.1        SnowballC_0.7.1
#>  [40] labeling_0.4.3         urltools_1.7.3.1       timechange_0.4.0
#>  [43] polyclip_1.10-7        httr_1.4.9             compiler_4.6.1
#>  [46] here_1.0.2             bit64_4.8.6            withr_3.0.3
#>  [49] S7_0.2.2               viridis_0.6.5          ggforce_0.5.0
#>  [52] MASS_7.3-65            bibliometrixData_0.3.0 tools_4.6.1
#>  [55] otel_0.2.0             stopwords_2.3          zip_3.0.2
#>  [58] httpuv_1.6.17          rentrez_1.2.4          promises_1.5.0
#>  [61] grid_4.6.1             stringdist_0.9.17      reshape2_1.4.5
#>  [64] generics_0.1.4         gtable_0.3.6           tzdb_0.5.0
#>  [67] rscopus_0.9.0          ca_0.72                data.table_1.18.6.1
#>  [70] hms_1.1.4              xml2_1.6.0             utf8_1.2.6
#>  [73] ggrepel_0.9.8          pillar_1.11.1          nsyllable_1.0.1
```

```
#>  [76] vroom_1.7.1            later_1.4.8            tweenr_2.0.3
#>  [79] lattice_0.22-9         bit_4.6.0              tidyselect_1.2.1
#>  [82] tm_0.7-20              miniUI_0.1.2           downlit_0.4.5
#>  [85] knitr_1.52             gridExtra_2.3.1        NLP_0.3-3
#>  [88] bookdown_0.48          stats4_4.6.1           crul_1.6.0
#>  [91] xfun_0.62              graphlayouts_1.2.5     matrixStats_1.5.0
#>  [94] DT_0.34.0              humaniformat_0.6.0     stringi_1.8.9
#>  [97] qpdf_1.4.1             yaml_2.3.12            evaluate_1.0.5
#> [100] httpcode_0.3.0         cli_3.6.6              xtable_1.8-8
#> [103] jquerylib_0.1.4        Rcpp_1.1.2             readxl_1.6.0
#> [106] triebeard_0.4.1        XML_3.99-0.25          parallel_4.6.1
#> [109] assertthat_0.2.1       pubmedR_1.0.2          slam_0.1-56
#> [112] viridisLite_0.4.3      scales_1.4.0           crayon_1.5.3
#> [115] openxlsx_4.2.9         rlang_1.3.0            fastmatch_1.1-8
```